

Class 12: RNA Seq Analysis

Dan Vu (PID: A17380158)

Table of contents

Background	1
Data import	1
Toy with differential gene expression	3
DESeq2 Analysis	7
Adding Annotation Data	9
Volcano Plot	11
Save our results	12
Add gene annotation	12

Background

Today we will analyze some RNASeq data from Himes et al. on the effects of a common steroid on airway smooth muscle cells (ASM cells).

Our starting point is the “counts” data and “metadata” that contain the count values for each gene in their different experiments (i.e. cell lines with or without the drug)

Data import

```
metadata <- read.csv("airway_metadata.csv")
counts <- read.csv("airway_scaledcounts.csv", row.names=1)
```

Let's see what these datasets look like

```
head(counts)
```

	SRR1039508	SRR1039509	SRR1039512	SRR1039513	SRR1039516
ENSG000000000003	723	486	904	445	1170
ENSG000000000005	0	0	0	0	0
ENSG000000000419	467	523	616	371	582
ENSG000000000457	347	258	364	237	318
ENSG000000000460	96	81	73	66	118
ENSG000000000938	0	0	1	0	2
	SRR1039517	SRR1039520	SRR1039521		
ENSG000000000003	1097	806	604		
ENSG000000000005	0	0	0		
ENSG000000000419	781	417	509		
ENSG000000000457	447	330	324		
ENSG000000000460	94	102	74		
ENSG000000000938	0	0	0		

```
head(metadata)
```

	id	dex	celltype	geo_id
1	SRR1039508	control	N61311	GSM1275862
2	SRR1039509	treated	N61311	GSM1275863
3	SRR1039512	control	N052611	GSM1275866
4	SRR1039513	treated	N052611	GSM1275867
5	SRR1039516	control	N080611	GSM1275870
6	SRR1039517	treated	N080611	GSM1275871

Q1. How many genes are in this dataset?

```
nrow(counts)
```

[1] 38694

Q2. How many different experiments are there? How many “control” cell lines are there?

```
nrow(metadata)
```

[1] 8

```
sum(metadata$dex == "control")
```

[1] 4

There are 4 experimental ‘control’ cell lines and 4 ‘control’ cell lines.

Toy with differential gene expression

To start our analysis let's calculate the mean counts for all genes in the "control" experiments.

1. Extract all "control" columns from the `counts` object.
 2. Calculate the mean for all rows(i.e. genes) of these "control" columns.
- 3-4. Do the same for "treated" 5. Compare these with `control.mean` values to see if theres any significant difference.

```
library(dplyr)
```

Attaching package: 'dplyr'

The following objects are masked from 'package:stats':

`filter`, `lag`

The following objects are masked from 'package:base':

`intersect`, `setdiff`, `setequal`, `union`

```
control inds <- metadata$dex == "control"  
control.counts <- counts[,control inds]
```

```
meanrows.control.counts <- rowMeans(control.counts)
```

Q3. How would you make the above code in either approach more robust? Is there a function that could help here?

Yes, the above code could be more robust with the function `rowMeans`.

Q4. Follow the same procedure for the treated samples (i.e. calculate the mean per gene across drug treated samples and assign to a labeled vector called `treated.mean`)

Now let's do the same for the experimental cells.

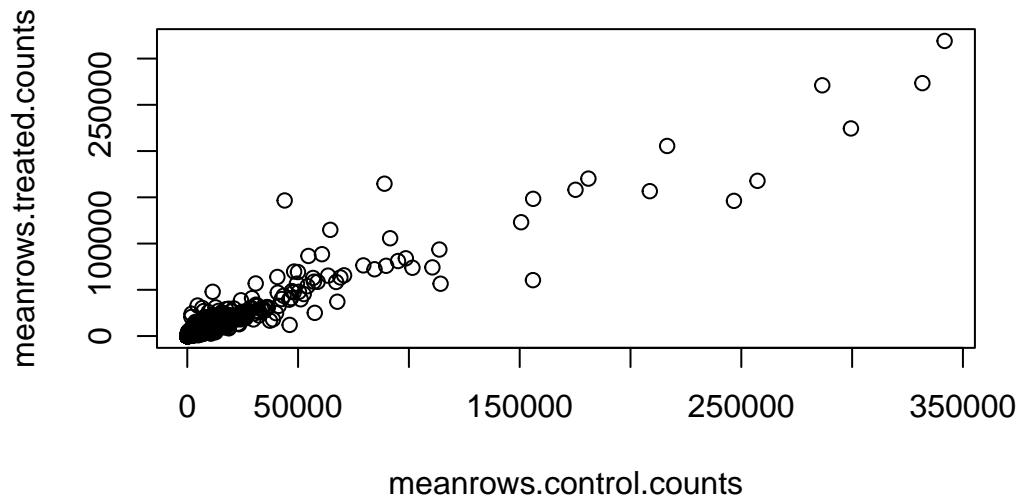
```
treated inds <- metadata$dex == "treated"  
treated.counts <- counts[,treated inds]
```

```
meanrows.treated.counts <- rowMeans(treated.counts)
```

```
meancounts <- data.frame(meanrows.control.counts, meanrows.treated.counts)
head(meancounts)
```

	meanrows.control.counts	meanrows.treated.counts
ENSG000000000003	900.75	658.00
ENSG000000000005	0.00	0.00
ENSG000000000419	520.50	546.00
ENSG000000000457	339.75	316.50
ENSG000000000460	97.25	78.75
ENSG000000000938	0.75	0.00

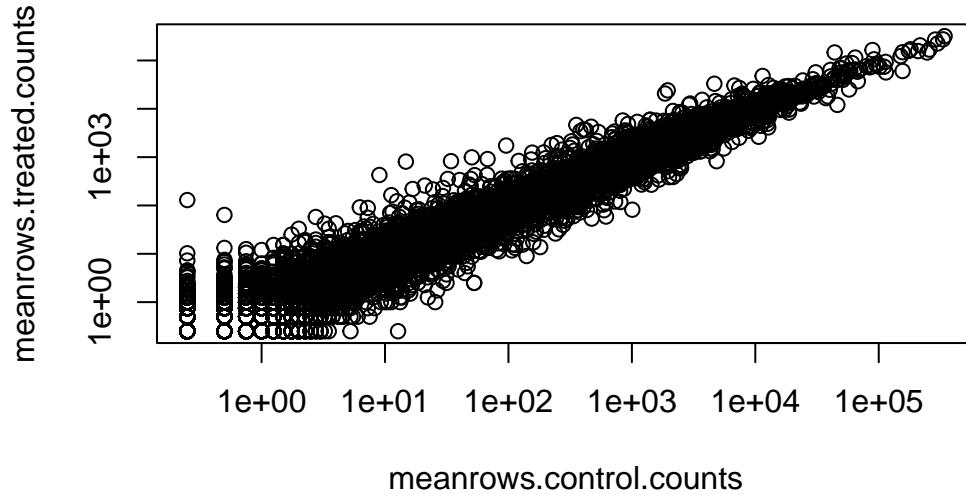
```
plot(meancounts)
```



```
plot(meancounts, log="xy")
```

Warning in xy.coords(x, y, xlabel, ylabel, log): 15032 x values <= 0 omitted
from logarithmic plot

```
Warning in xy.coords(x, y, xlabel, ylabel, log): 15281 y values <= 0 omitted  
from logarithmic plot
```



Q5. You could also use the ggplot2 package to make this figure producing the plot below. What geom_?() function would you use for this plot?

I would use the `geom_point()` function.

Q6. Try plotting both axes on a log scale. What is the argument to `plot()` that allows you to do this?

The argument is `plot(meancounts, log="xy")`.

We often talk metrics like “log₂ fold-change”

```
# treated/control  
log2(10/10)
```

```
[1] 0
```

```
log2(10/40)
```

```
[1] -2
```

Let's calculate the log2 fold change for our treated over control mean counts.

```
meancounts$log2fc <-  
log2(meancounts$meanrows.treated.counts /  
     meancounts$meanrows.control.counts)  
  
head(meancounts)  
  
          meanrows.control.counts meanrows.treated.counts      log2fc  
ENSG000000000003             900.75            658.00 -0.45303916  
ENSG000000000005              0.00             0.00        NaN  
ENSG000000000419             520.50            546.00  0.06900279  
ENSG000000000457             339.75            316.50 -0.10226805  
ENSG000000000460             97.25             78.75 -0.30441833  
ENSG000000000938              0.75             0.00       -Inf  
  
zero.vals <- which(meancounts[,1:2]==0, arr.ind=TRUE)  
  
to.rm <- unique(zero.vals[,1])  
mycounts <- meancounts[-to.rm,]  
head(mycounts)  
  
          meanrows.control.counts meanrows.treated.counts      log2fc  
ENSG000000000003             900.75            658.00 -0.45303916  
ENSG000000000419             520.50            546.00  0.06900279  
ENSG000000000457             339.75            316.50 -0.10226805  
ENSG000000000460             97.25             78.75 -0.30441833  
ENSG000000000971            5219.00            6687.50  0.35769358  
ENSG000000001036            2327.00            1785.75 -0.38194109
```

Q7. What is the purpose of the arr.ind argument in the which() function call above? Why would we then take the first column of the output and need to call the unique() function?

The purpose of the arr.ind argument in the which() function is to report the positions where there are any “TRUE” values, allowing us to ignore any zero counts.

A common “rule of thumb” is a log2 fold change cutoff of +2 or -2 to call genes “Up regulated” or “Down regulated”.

Q8. Using the up.ind vector above can you determine how many up regulated genes we have at the greater than 2 fc level?

This is the number of up-regulated genes.

```
sum(meancounts$log2fc >= +2, na.rm = T)
```

```
[1] 1910
```

Q9. Using the down.ind vector above can you determine how many down regulated genes we have at the greater than 2 fc level?

This is the number of down-regulated genes.

```
sum(meancounts$log2fc <= -2, na.rm = T)
```

```
[1] 2330
```

Q10. Do you trust these results? Why or why not?

Although these fold changes are large, whether or not they are considered statistically significant between the two types of cell lines is still to be determined, so these results cannot be trusted.

DESeq2 Analysis

```
library(DESeq2)
```

For DESeq2 analysis we need three things. - count values (countData) - metadata telling us about the columns in countData (colData) - design of the experiment (i.e. what do you want to compare)

Our first function from DESeq2 will setup the input required for analysis by storing all these 3 things together.

```
dds <- DESeqDataSetFromMatrix(countData = counts,
                                colData = metadata,
                                design = ~dex)
```

```
converting counts to integer mode
```

```
Warning in DESeqDataSet(se, design = design, ignoreRank): some variables in
design formula are characters, converting to factors
```

The main function in DESeq2 that runs the analysis is called `DESeq()`

```
dds <- DESeq(dds)
```

estimating size factors

estimating dispersions

gene-wise dispersion estimates

mean-dispersion relationship

final dispersion estimates

fitting model and testing

```
res <- results(dds)
head(res)
```

log2 fold change (MLE): dex treated vs control

Wald test p-value: dex treated vs control

DataFrame with 6 rows and 6 columns

	baseMean	log2FoldChange	lfcSE	stat	pvalue
	<numeric>	<numeric>	<numeric>	<numeric>	<numeric>
ENSG000000000003	747.194195	-0.350703	0.168242	-2.084514	0.0371134
ENSG000000000005	0.000000	NA	NA	NA	NA
ENSG000000000419	520.134160	0.206107	0.101042	2.039828	0.0413675
ENSG000000000457	322.664844	0.024527	0.145134	0.168996	0.8658000
ENSG000000000460	87.682625	-0.147143	0.256995	-0.572550	0.5669497
ENSG000000000938	0.319167	-1.732289	3.493601	-0.495846	0.6200029

	padj
	<numeric>
ENSG000000000003	0.163017
ENSG000000000005	NA
ENSG000000000419	0.175937
ENSG000000000457	0.961682
ENSG000000000460	0.815805
ENSG000000000938	NA

Adding Annotation Data

```
library("AnnotationDbi")
```

Attaching package: 'AnnotationDbi'

The following object is masked from 'package:dplyr':

```
select
```

```
library("org.Hs.eg.db")
```

```
columns(org.Hs.eg.db)
```

```
[1] "ACNUM"      "ALIAS"       "ENSEMBL"     "ENSEMLPROT"  "ENSEMLTRANS"
[6] "ENTREZID"   "ENZYME"      "EVIDENCE"    "EVIDENCEALL" "GENENAME"
[11] "GENETYPE"   "GO"          "GOALL"       "IPI"         "MAP"
[16] "OMIM"        "ONTOLOGY"    "ONTOLOGYALL" "PATH"        "PFAM"
[21] "PMID"        "PROSITE"     "REFSEQ"      "SYMBOL"     "UCSCKG"
[26] "UNIPROT"
```

```
res$symbol <- mapIds(org.Hs.eg.db,
                      keys=row.names(res), # Our genenames
                      keytype="ENSEMBL",      # The format of our genenames
                      column="SYMBOL",        # The new format we want to add
                      multiVals="first")
```

'select()' returned 1:many mapping between keys and columns

```
head(res)
```

```
log2 fold change (MLE): dex treated vs control
Wald test p-value: dex treated vs control
DataFrame with 6 rows and 7 columns
```

	baseMean	log2FoldChange	lfcSE	stat	pvalue
	<numeric>	<numeric>	<numeric>	<numeric>	<numeric>
ENSG000000000003	747.194195	-0.350703	0.168242	-2.084514	0.0371134
ENSG000000000005	0.000000	NA	NA	NA	NA
ENSG000000000419	520.134160	0.206107	0.101042	2.039828	0.0413675
ENSG00000000457	322.664844	0.024527	0.145134	0.168996	0.8658000
ENSG00000000460	87.682625	-0.147143	0.256995	-0.572550	0.5669497
ENSG00000000938	0.319167	-1.732289	3.493601	-0.495846	0.6200029
	padj	symbol			
	<numeric>	<character>			
ENSG000000000003	0.163017	TSPAN6			
ENSG000000000005	NA	TNMD			
ENSG000000000419	0.175937	DPM1			
ENSG00000000457	0.961682	SCYL3			
ENSG00000000460	0.815805	FIRRM			
ENSG00000000938	NA	FGR			

Q11. Run the mapIds() function two more times to add the Entrez ID and UniProt accession and GENENAME as new columns called res\$entrez, res\$uniprot and res\$genename.

```
res$entrez <- mapIds(org.Hs.eg.db,
                      keys=row.names(res),
                      column="ENTREZID",
                      keytype="ENSEMBL",
                      multiVals="first")
```

'select()' returned 1:many mapping between keys and columns

```
res$uniprot <- mapIds(org.Hs.eg.db,
                      keys=row.names(res),
                      column="UNIPROT",
                      keytype="ENSEMBL",
                      multiVals="first")
```

'select()' returned 1:many mapping between keys and columns

```
res$genename <- mapIds(org.Hs.eg.db,
                      keys=row.names(res),
                      column="GENENAME",
                      keytype="ENSEMBL",
                      multiVals="first")
```

```
'select()' returned 1:many mapping between keys and columns
```

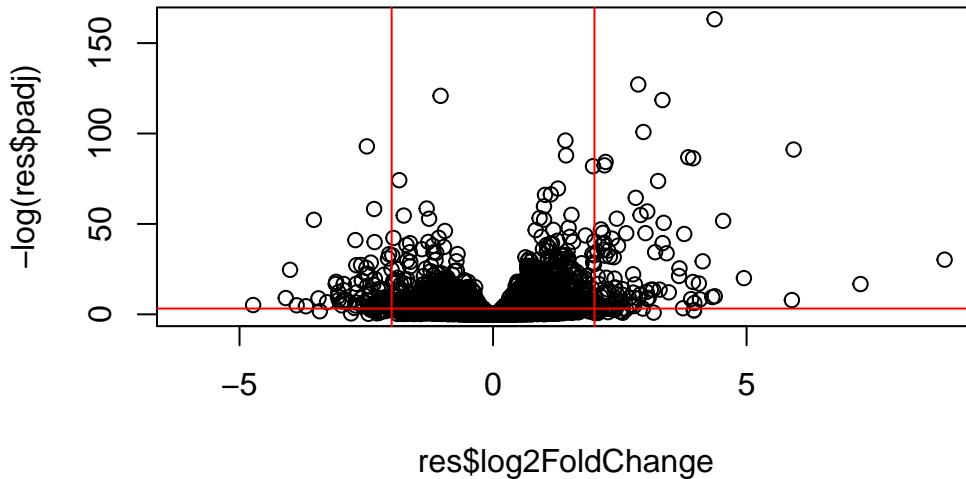
```
head(res)
```

```
log2 fold change (MLE): dex treated vs control
Wald test p-value: dex treated vs control
DataFrame with 6 rows and 10 columns
  baseMean log2FoldChange    lfcSE      stat     pvalue
  <numeric>      <numeric> <numeric> <numeric> <numeric>
ENSG000000000003 747.194195 -0.350703  0.168242 -2.084514 0.0371134
ENSG000000000005  0.000000      NA       NA       NA       NA
ENSG000000000419 520.134160  0.206107  0.101042  2.039828 0.0413675
ENSG000000000457 322.664844  0.024527  0.145134  0.168996 0.8658000
ENSG000000000460 87.682625 -0.147143  0.256995 -0.572550 0.5669497
ENSG000000000938 0.319167 -1.732289  3.493601 -0.495846 0.6200029
  padj      symbol     entrez     uniprot
  <numeric> <character> <character> <character>
ENSG000000000003 0.163017   TSPAN6      7105 AOA087WYV6
ENSG000000000005  NA        TNMD       64102 Q9H2S6
ENSG000000000419 0.175937   DPM1       8813 H0Y368
ENSG000000000457 0.961682   SCYL3      57147 X6RHX1
ENSG000000000460 0.815805   FIRRM      55732 A6NFP1
ENSG000000000938  NA        FGR        2268 B7Z6W7
  genename
  <character>
ENSG000000000003      tetraspanin 6
ENSG000000000005      tenomodulin
ENSG000000000419 dolichyl-phosphate m..
ENSG000000000457 SCY1 like pseudokina..
ENSG000000000460 FIGNL1 interacting r..
ENSG000000000938 FGR proto-oncogene, ..
```

Volcano Plot

This is a common summary result figure from these types of experiments and plot the log2 fold-change vs the p-value

```
plot(res$log2FoldChange, -log(res$padj))
abline(v=c(-2,2), col="red")
abline(h=-log(0.04), col="red")
```



Save our results

```
write.csv(res, file="my_results.csv")
```

Add gene annotation

To help make sense of our results and communicate them to other folks we need to add some more annotation to our main `res` object.

We will use two packages to first map IDs to different formats including the classic gene “symbol” gene name.

```
library(AnnotationDbi)
library(org.Hs.eg.db)
```

Let’s see what is in the second package and what exactly it can do for us.

```
columns(org.Hs.eg.db)
```

```
[1] "ACCCNUM"      "ALIAS"        "ENSEMBL"       "ENSEMBLPROT"   "ENSEMBLTRANS"
[6] "ENTREZID"     "ENZYME"       "EVIDENCE"      "EVIDENCEALL"   "GENENAME"
[11] "GENETYPE"     "GO"          "GOALL"         "IPI"          "MAP"
[16] "OMIM"          "ONTOLOGY"     "ONTOLOGYALL"  "PATH"         "PFAM"
[21] "PMID"          "PROSITE"      "REFSEQ"        "SYMBOL"       "UCSCKG"
[26] "UNIPROT"
```

We can translate or “map” IDs between any of these 26 databases using the `mapIDs()` function.

```
res$symbol <- mapIDs(keys = row.names(res),
  keytype = "ENSEMBL",
  x = org.Hs.eg.db,
  column = "SYMBOL"
)
```

```
'select()' returned 1:many mapping between keys and columns
```

```
head(res)
```

```
log2 fold change (MLE): dex treated vs control
Wald test p-value: dex treated vs control
DataFrame with 6 rows and 10 columns
  baseMean log2FoldChange      lfcSE      stat    pvalue
  <numeric>      <numeric> <numeric> <numeric> <numeric>
ENSG000000000003 747.194195 -0.350703  0.168242 -2.084514 0.0371134
ENSG000000000005  0.000000      NA        NA        NA        NA
ENSG00000000419  520.134160  0.206107  0.101042  2.039828 0.0413675
ENSG00000000457  322.664844  0.024527  0.145134  0.168996 0.8658000
ENSG00000000460  87.682625 -0.147143  0.256995 -0.572550 0.5669497
ENSG00000000938  0.319167 -1.732289  3.493601 -0.495846 0.6200029
  padj      symbol      entrez      uniprot
  <numeric> <character> <character> <character>
ENSG000000000003 0.163017    TSPAN6      7105    AOA087WYV6
ENSG000000000005  NA        TNMD       64102    Q9H2S6
ENSG00000000419  0.175937    DPM1       8813    HOY368
ENSG00000000457  0.961682    SCYL3      57147    X6RHX1
ENSG00000000460  0.815805    FIRRM      55732    A6NFP1
ENSG00000000938  NA        FGR        2268    B7Z6W7
  genename
  <character>
```

```

ENSG000000000003      tetraspanin 6
ENSG000000000005      tenomodulin
ENSG000000000419 dolichyl-phosphate m..
ENSG000000000457 SCY1 like pseudokina..
ENSG000000000460 FIGNL1 interacting r..
ENSG000000000938 FGR proto-oncogene, ..

```

Add the mappings for “GENENAME” and “ENTREZID” and sore as `res$genename` and `res$entrez`.

```

res$genename <- mapIds(keys = row.names(res),
  keytype = "ENSEMBL",
  x = org.Hs.eg.db,
  column = "GENENAME"
)

```

```
'select()' returned 1:many mapping between keys and columns
```

```
head(res)
```

```

log2 fold change (MLE): dex treated vs control
Wald test p-value: dex treated vs control
DataFrame with 6 rows and 10 columns
  baseMean log2FoldChange      lfcSE      stat     pvalue
  <numeric>      <numeric> <numeric> <numeric> <numeric>
ENSG000000000003 747.194195 -0.350703  0.168242 -2.084514 0.0371134
ENSG000000000005  0.000000       NA        NA        NA        NA
ENSG000000000419 520.134160  0.206107  0.101042  2.039828 0.0413675
ENSG000000000457 322.664844  0.024527  0.145134  0.168996 0.8658000
ENSG000000000460 87.682625  -0.147143  0.256995 -0.572550 0.5669497
ENSG000000000938 0.319167  -1.732289  3.493601 -0.495846 0.6200029
  padj      symbol      entrez      uniprot
  <numeric> <character> <character> <character>
ENSG000000000003 0.163017    TSPAN6      7105 AOA087WYV6
ENSG000000000005   NA        TNMD      64102 Q9H2S6
ENSG000000000419 0.175937    DPM1      8813 H0Y368
ENSG000000000457 0.961682    SCYL3      57147 X6RHX1
ENSG000000000460 0.815805    FIRRM      55732 A6NFP1
ENSG000000000938   NA        FGR       2268 B7Z6W7
  genename
  <character>

```

```

ENSG000000000003      tetraspanin 6
ENSG000000000005      tenomodulin
ENSG000000000419 dolichyl-phosphate m..
ENSG000000000457 SCY1 like pseudokina..
ENSG000000000460 FIGNL1 interacting r..
ENSG000000000938 FGR proto-oncogene, ..

```

```

res$entrezid <- mapIds(keys = row.names(res),
  keytype = "ENSEMBL",
  x = org.Hs.eg.db,
  column = "ENTREZID"
)

```

```
'select()' returned 1:many mapping between keys and columns
```

```
head(res)
```

```

log2 fold change (MLE): dex treated vs control
Wald test p-value: dex treated vs control
DataFrame with 6 rows and 11 columns
  baseMean log2FoldChange      lfcSE      stat     pvalue
  <numeric>      <numeric> <numeric> <numeric> <numeric>
ENSG000000000003 747.194195 -0.350703  0.168242 -2.084514 0.0371134
ENSG000000000005  0.000000       NA        NA        NA        NA
ENSG000000000419 520.134160  0.206107  0.101042  2.039828 0.0413675
ENSG000000000457 322.664844  0.024527  0.145134  0.168996 0.8658000
ENSG000000000460 87.682625 -0.147143  0.256995 -0.572550 0.5669497
ENSG000000000938 0.319167 -1.732289  3.493601 -0.495846 0.6200029
  padj      symbol      entrez      uniprot
  <numeric> <character> <character> <character>
ENSG000000000003 0.163017    TSPAN6      7105 AOA087WYV6
ENSG000000000005   NA        TNMD      64102 Q9H2S6
ENSG000000000419 0.175937    DPM1       8813 H0Y368
ENSG000000000457 0.961682    SCYL3      57147 X6RHX1
ENSG000000000460 0.815805    FIRRM      55732 A6NFP1
ENSG000000000938   NA        FGR       2268 B7Z6W7
  genename      entrezid
  <character> <character>
ENSG000000000003      tetraspanin 6      7105
ENSG000000000005      tenomodulin      64102
ENSG000000000419 dolichyl-phosphate m..      8813

```

```
ENSG00000000457 SCY1 like pseudokina..      57147
ENSG00000000460 FIGNL1 interacting r..      55732
ENSG00000000938 FGR proto-oncogene, ..     2268
```

##Pathway Analysis

There are lots of bioconductor packages to do this type of analysis.

For now let's try one called **gage**.

```
library(gage)
```

```
library(gageData)
library(pathview)
```

```
#####
# Pathview is an open source software package distributed under GNU General
# Public License version 3 (GPLv3). Details of GPLv3 is available at
# http://www.gnu.org/licenses/gpl-3.0.html. Particullary, users are required to
# formally cite the original Pathview paper (not just mention it) in publications
# or products. For details, do citation("pathview") within R.
```

The pathview downloads and uses KEGG data. Non-academic uses may require a KEGG license agreement (details at <http://www.kegg.jp/kegg/legal.html>).

```
#####
```

To use **gage** I need two things.

- a names vector of fold-change values for our DEGs (our geneset of interest)
- a set of pathways or genesets to use for annotation

```
c("barry" = 5, "lisa" = 10)
```

```
barry  lisa
5      10
```

```
foldchanges <- res$log2FoldChange
names(foldchanges) <- res$symbol
head(foldchanges)
```

```

TSPAN6          TNMD          DPM1          SCYL3          FIRRM          FGR
-0.35070296    NA   0.20610728  0.02452701 -0.14714263 -1.73228897

```

```

data("kegg.sets.hs")
keggres = gage(foldchanges, gsets=kegg.sets.hs)

```

In our results object we have:

```
attributes(keggres)
```

```

$names
[1] "greater" "less"     "stats"

```

```
head(keggres$less, 3)
```

	p.geomean	stat.mean	p.val	q.val
hsa00232 Caffeine metabolism	NA	NaN	NA	NA
hsa00983 Drug metabolism - other enzymes	NA	NaN	NA	NA
hsa01100 Metabolic pathways	NA	NaN	NA	NA
	set.size	exp1		
hsa00232 Caffeine metabolism	0	NA		
hsa00983 Drug metabolism - other enzymes	0	NA		
hsa01100 Metabolic pathways	0	NA		

Let's look at one of these pathways (hsa05310 Asthma) with our genes colored up so we can see the overlap.

```
pathview(pathway.id = "hsa05130", gene.data = foldchanges)
```

```

Warning: None of the genes or compounds mapped to the pathway!
Argument gene.idtype or cpd.idtype may be wrong.

```

```
'select()' returned 1:1 mapping between keys and columns
```

```
Info: Working in directory /Users/danvu/Desktop/BIMM_143/Class_12
```

```
Info: Writing image file hsa05130.pathview.png
```

Add this pathway figure to our lab report.

